

1. Problem & Target User

Migration decisions involve multiple stakeholders: engineering (performance, latency), product (UX, velocity), finance (ROI), compliance (data governance), and architecture (vendor risk). They stall most often at evaluation, before any other stakeholder gets involved.

Problem selected: "Will this open model be good enough for my specific use case and prompts?" Cost and latency advantages speak for themselves once quality parity is demonstrated. Technical dependency migration is out of scope, they're too unique per customer to generalize.

Target user: Software/ML engineers at AI-forward startups or enterprise engineering orgs.

2. Goals & Metrics

Give engineers evidence that an open model matches their quality bar at lower cost so they can make a confident internal case for migration.

Metrics: model recommendation satisfaction (user rating), usage completion rate (% downloading summary), conversion rate (% contacting sales or moving to contract).

3. What I Built

A model recommender and evaluation tool that takes a user's existing prompt, recommends top Fireworks models for that use case, runs the prompt live, and compares latency, token usage, and cost side-by-side.

The call to action is downloading the comparison summary as an internal pitch artifact.

4. Key Choices

- **Evaluation phase only.** Solving the gateway question "is this good enough?" and unblocks downstream conversation. Full migration tooling is too unique per customer to generalize.
- **Recommendation before comparison.** The tool suggests top Fireworks models based on use case type rather than asking users to choose blind. Mirrors how Fireworks sells to customers already having a solution and makes the recommendation defensible.
- **Cost shown after the run.** Actual token counts from the live API response rather than upfront estimates makes cost comparison reliable
- **User satisfaction feedback loop.** The tool includes an in-line quality rating after each model output — users indicate whether the response met their bar, which captures recommendation accuracy signal without adding friction. This feedback loop ensures that the users leave with the highest satisfaction.
- **Call to action.** A shareable summary the engineer brings to their manager drives sales conversations organically at the right moment rather than forcing premature conversion.

5. Risks & Considerations

- Single prompt is a snapshot only, latency and cost at scale may differ significantly due to workflow
- LLM only, does not cover multimodal, agentic, or multi-model routing use cases
- Accountability gap when an open model underperforms in production
- Enterprise deals may require hands-on customer success a self-serve tool cannot replace

6. What's Next

- Batch prompt testing: 3-5 prompts for a statistically meaningful confidence score. Medium effort, high credibility impact.
- Use case classifier: auto-detect prompt type and refine model recommendations per category. Medium effort improves recommendation quality directly.